

Explanation of Ride Pricing Model Implementation

0.1 Overview

This document describes a Python-based pipeline for processing ride-sharing data and modeling ride prices to derive coefficients for a dynamic pricing algorithm. It improves upon an earlier implementation (`linearRegression.py`) that predicted ride durations using uniformly distributed data from `generate_ride_data.py` which did not reflect real-world congestion patterns. The updated pipeline, implemented in `PriceTargetLinearRegression.py`, uses surge-based data from `SurgePricingDataGeneration.py` with demand peaks at 8 AM and 5 PM, targeting ride prices in South African Rand (ZAR) to support dynamic pricing.

0.2 Library Imports

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.linear_model import LinearRegression
4 from sklearn.linear_model import RidgeCV, LassoCV, ElasticNetCV
5 from datetime import timedelta
6 from sqlalchemy import create_engine
7 from dotenv import load_dotenv
8 import os
9 import urllib
```

The pipeline uses:

- `pandas, numpy`: For data manipulation and numerical computations.
- `sklearn.linear_model`: For fitting regression models to predict prices.
- `datetime.timedelta`: For calculating time windows in hourly statistics.
- `sqlalchemy, urllib`: For SQL Server database connectivity.
- `dotenv, os`: For secure loading of database credentials from a `.env` file.

0.3 Environment and Database Configuration

```
1 load_dotenv()
2 DB_HOST = os.getenv('SQLSERVER_HOST')
3 DB_PORT = os.getenv('SQLSERVER_PORT', '1433')
4 DB_NAME = os.getenv('SQLSERVER_DB')
5 DB_USER = os.getenv('SQLSERVER_USER')
6 DB_PASSWORD = os.getenv('SQLSERVER_PASSWORD')
7 odbc_str = (
8     f"DRIVER=ODBC_Driver_17_for_SQL_Server;"
9     f"SERVER={DB_HOST},{DB_PORT};"
10    f"DATABASE={DB_NAME};"
11    f"UID={DB_USER};"
```

```

12     f"PWD={DB_PASSWORD};"
13     "TrustServerCertificate=yes;"
14     "ConnectionTimeout=30;"
15     "CommandTimeout=60;"
16 )
17 connection_uri = f"mssql+pyodbc:///odbc_connect={urllib.parse.
    quote_plus(odbc_str)}"
18 engine = create_engine(
19     connection_uri,
20     pool_timeout=30,
21     pool_recycle=300,
22     pool_pre_ping=True
23 )

```

Environment variables are loaded from a `.env` file for secure database credentials. An ODBC connection string with timeout settings ensures robust connectivity to a SQL Server database via sqlalchemy's `create_engine`.

0.4 Data Extraction

```

1 df_requests = pd.read_sql("SELECT_*_FROM_ride_requests", con=engine)
2 df_completion = pd.read_sql("SELECT_*_FROM_ride_completion_delay",
    con=engine)
3 df_acceptance = pd.read_sql("SELECT_*_FROM_ride_acceptance_delay",
    con=engine)
4 df_matches = pd.read_sql("SELECT_*_FROM_ride_matches", con=engine)

```

Four datasets are extracted from SQL Server:

- `ride_requests`: Contains ride request details (e.g., `ride_id`, `request_time`, `pickup_lat`).
- `ride_acceptance_delay`: Records driver acceptance times and delays.
- `ride_completion_delay`: Includes ride durations and completion times.
- `ride_matches`: Stores ride matches with price and status information.

Unlike the previous version, which used a precomputed `ride_pricing_features` table, this pipeline processes raw ride data directly, enabling dynamic feature generation.

0.5 Data Cleaning

```

1 for df in [df_requests, df_completion, df_acceptance, df_matches]:
2     df.columns = df.columns.str.strip().str.replace('""', '').str.
        replace("'", '')
3     for col in df.columns:
4         if df[col].dtype == 'object':
5             df[col] = df[col].astype(str).str.replace(',', '.')

```

```

6         try:
7             df[col] = pd.to_numeric(df[col])
8         except:
9             pass

```

Column names are cleaned to remove extraneous quotes, and comma-based decimal separators are converted to dots to ensure numerical consistency.

0.6 Datetime Parsing

```

1 datetime_cols = {
2     'request_time': [df_requests, df_acceptance],
3     'pickup_time': [df_completion],
4     'accepted_time': [df_acceptance],
5     'completed_at': [df_completion, df_matches],
6     'matched_at': [df_matches],
7     'driver_response_at': [df_matches],
8     'started_at': [df_matches],
9     'arrived_at': [df_matches]
10 }
11 for col_name, dataframes in datetime_cols.items():
12     for df in dataframes:
13         if col_name in df.columns:
14             df[col_name] = pd.to_datetime(df[col_name], format='%Y-%
                m-%d_%H:%M:%S.%f', errors='coerce')

```

Datetime fields are parsed with a specific format to ensure accurate temporal analysis, with invalid entries coerced to handle errors.

0.7 Feature Engineering

The pipeline constructs a feature dataset from `ride_matches`, adding:

- **Distance:** Euclidean distance (`distance_km`) calculated as:

$$\text{distance_km} = \sqrt{(\text{dropoff_lat} - \text{pickup_lat})^2 + (\text{dropoff_lng} - \text{pickup_lng})^2} \times 111 \quad (1)$$

- **Temporal Features:** `request_hour` and `request_day` extracted from `matched_at`.
- **Acceptance Time:** `current_accept_time` merged from `ride_acceptance_delay`, with missing values filled by the median.

0.8 Hourly Statistics Computation

```

1 def compute_hourly_stats(df_features, df_requests, df_acceptance,
2     df_completion):
3     stats = []
4     base_date = (df_features['matched_at'].iloc[0].floor('D'))

```

```

4         if not df_features['matched_at'].isna().all()
5         else pd.to_datetime('2025-01-01'))
6     for idx, row in df_features.iterrows():
7         ride_id = row['ride_id']
8         hour = row['request_hour']
9         lag_hour = max(0, hour - 1)
10        lead_hour = min(23, hour + 1)
11        prev_week_start = base_date - timedelta(weeks=1) + timedelta(
12            hours=lag_hour)
13        prev_week_end = base_date - timedelta(weeks=1) + timedelta(
14            hours=lead_hour + 1)
15        mask_requests = (
16            (df_requests['request_time'] >= prev_week_start) &
17            (df_requests['request_time'] < prev_week_end) &
18            (df_requests['request_time'].dt.hour.between(lag_hour,
19                lead_hour))
20        )
21        request_count = df_requests[mask_requests].shape[0]
22        hourly_counts = df_requests[mask_requests].groupby(
23            df_requests['request_time'].dt.hour).size()
24        avg_request_count = hourly_counts.mean() if not
25            hourly_counts.empty else 0
26        stdev_request_count = hourly_counts.std() if len(
27            hourly_counts) > 1 else 0
28        # ... (similar for acceptance and completion)
29        z_request_count = (request_count - avg_request_count) /
30            stdev_request_count if stdev_request_count > 0 else 0
31        z_accept_time = (row['current_accept_time'] -
32            avg_accept_time) / stdev_accept_time if stdev_accept_time
33            > 0 else 0
34        stats.append({
35            'ride_id': ride_id,
36            'request_count': request_count,
37            'avg_request_count': avg_request_count,
38            'stdev_request_count': stdev_request_count,
39            'avg_accept_time': avg_accept_time,
40            'stdev_accept_time': stdev_accept_time,
41            'avg_ride_duration': avg_ride_duration,
42            'stdev_ride_duration': stdev_ride_duration,
43            'z_request_count': z_request_count,
44            'z_accept_time': z_accept_time
45        })
46    return pd.DataFrame(stats)

```

The `compute_hourly_stats` function calculates historical statistics for the previous week's data within a 3-hour window (current hour ± 1), including:

- Request counts, averages, and standard deviations.
- Average and standard deviation of acceptance times and ride durations.

- Z-scores for request counts and acceptance times, calculated as:

$$z = \frac{\text{current} - \text{mean}}{\text{std_dev}} \quad (2)$$

Rows with zero statistics are filtered out to ensure meaningful data.

0.9 Merging and Filtering

Computed statistics are merged with the feature dataset:

```
1 df_merged = df_features.merge(stats_df, on='ride_id', how='left')
2 df_merged = df_merged[df_merged[stat_columns].ne(0).any(axis=1)]
```

This ensures only rows with non-zero statistics are used for modeling.

0.10 Model Fitting

Four regression models are trained to predict ride prices (ZAR), replacing the previous duration-based approach:

- **Linear Regression:** Fits a linear model, minimizing squared errors. Sensitive to outliers but interpretable.
- **Ridge Regression:** Uses L2 regularization to mitigate overfitting, with cross-validated alpha selection.
- **Lasso Regression:** Applies L1 regularization for feature selection, potentially zeroing coefficients, with cross-validated alpha.
- **ElasticNet Regression:** Combines L1 and L2 penalties for balanced regularization, with cross-validated parameters.

Features used:

- distance_km, current_accept_time, request_count
- avg_accept_time, avg_ride_duration
- z_request_count, z_accept_time

Models are evaluated using Mean Squared Error (MSE) and R² scores, with coefficients interpreted for pricing impact (e.g., price increase per kilometer).

0.11 Generating Pricing Coefficients

```
1 feature_to_pricing = {
2     'distance_km': 'baseRate',
3     'request_count': 'coeffRequests',
4     'avg_accept_time': 'coeffAcceptTime',
```

```

5     'avg_ride_duration': 'coeffRideDuration',
6     'z_request_count': 'stdDevFactor',
7     'z_accept_time': 'stdDevFactor'
8 }
9 for name, coefs in model_results.items():
10     included = [abs(c) for f, c in coefs.items() if f != '
11         distance_km']
12     total = sum(included) or 1
13     std_dev_features = ['z_request_count', 'z_accept_time']
14     std_dev_sum = sum(abs(coefs[f]) for f in std_dev_features) / len
15         (std_dev_features) if std_dev_features else 0
16     for f in feature_cols:
17         key = feature_to_pricing.get(f)
18         if key == 'baseRate':
19             print(f"DECLARE_{key}_FLOAT={linear_model.intercept_
20                 :.2f};--Base_price_from_intercept")
21         elif key == 'stdDevFactor':
22             continue
23         else:
24             coef_val = coefs[f]
25             print(f"DECLARE_{key}_FLOAT={coef_val:.4f};")
26 if std_dev_sum > 0:
27     print(f"DECLARE_stdDevFactor_FLOAT={std_dev_sum:.4f};")

```

Features are mapped to pricing parameters, with non-distance coefficients normalized:

$$\text{norm_coef} = \frac{|\text{coef}|}{\text{total_sum}} \quad (3)$$

The stdDevFactor is computed as the mean of normalized z-score coefficients. Output is formatted for SQL stored procedures, e.g.:

```

1 DECLARE @coeffRequests FLOAT = 0.2300;
2 DECLARE @stdDevFactor FLOAT = 0.1400;

```

0.12 Data Generation

The pipeline uses surge-based data from SurgePricingDataGeneration.py, addressing the limitations of generate_ride_data.py's uniform distribution. Key features:

- **Demand Modeling:** Simulates demand peaks at 8 AM and 5 PM using Gaussian-like functions:

$$\text{demand} = 10 + 30 \exp\left(-\frac{(\text{hours} - 8)^2}{10}\right) + 40 \exp\left(-\frac{(\text{hours} - 17)^2}{12}\right) + \text{random noise} \quad (4)$$

- **Surge Pricing:** Calculates prices based on a demand-to-supply ratio, clipped between 1x and 3x, applied to a base ride price.

- **Variable Rides:** Scales ride volume by demand, with surge-dependent acceptance delays.

Data is generated for a single day, centered in Sandton, Gauteng, South Africa.

0.13 Conclusion

The updated pipeline leverages surge-based data, dynamic feature engineering, and price prediction to generate coefficients for a dynamic pricing algorithm. By addressing the uniform distribution limitation, it produces realistic, data-driven pricing coefficients suitable for integration into a stored procedure like `CalculateDynamicPrice`.